

# Benchopt: How standard is the performance of ResNet18 on CIFAR10?

Thomas Moreau INRIA Saclay



**CIFAR10:** Image classification of  $32 \times 32$  images from 10 classes.


**ResNet18:** *de facto* baseline for this task for many developments:

- ▶ Optimizers,
- ▶ Data Augmentation,
- ▶ Architecture,
- ▶ ...

⇒ However, replicating SOTA for this architecture  
is actually quite hard!

and many failed...

The goal of **Benchopt** is to make this step as easy as possible.



Best Adam

SGD + data aug. + momentum

Lookahead

Vanilla SGD

SGD + data aug. + momentum + cosine LR sched.

SGD + data aug.

Best SGD

Best SGD (TF/Keras)

## Extra comparisons

- ▶ Multiple datasets: MNIST, SVHN, CIFAR10, CIFAR100, ...
- ▶ Matching PyTorch and Tensorflow.

Component	PyTorch	TensorFlow/Keras
Bias in convolutions	<b>×</b>	<b>✓</b>
Decoupled weight decay scaling	<b>Multiplied by learning rate</b>	Completely decoupled
Batch normalization momentum	<b>0.9</b>	0.99
Conv2D weights init.	<b>Fan out, normal</b>	Fan average, uniform
Classification head init. (weights)	<b>Fan in, uniform</b>	Fan average, uniform
Classification head init. (bias)	<b>Fan in, uniform</b>	Zeros
Striding in convolutions	<b>Starts one off</b>	Ends one off
Variance estimation in batch norm	unbiased (eval)/biased (training)	biased

- ▶ Evaluated the impact of cross-validation,

# Benchmarks


You can run a similar benchmark with the correct baseline from:

[https://github.com/benchopt/benchmark\\_resnet\\_classif](https://github.com/benchopt/benchmark_resnet_classif)

Most of the heavy lifting have been done together with:



Zaccharie Ramzi



Pierre Ablin

Arxiv paper: <https://arxiv.org/abs/2206.13424>