

Robustifying `concurrent.futures`

Thomas Moreau - Olivier Grisel



école
normale
supérieure
paris—saclay



Embarassingly parallel computation in `python` using a pool of workers

Three API available:

- `multiprocessing`: first implementation.
- `concurrent.futures`: reimplementation using `multiprocessing` under the hood.
- `loky`: robustification of `concurrent.futures`.

The `concurrent.futures` API

The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor  
  
def fit_model(params):  
    # Heavy computation  
    return model
```

`fit_model` is the function that we want to run asynchronously.

The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:
```

`fit_model` is the function that we want to run asynchronously.

We instantiate a `ThreadPoolExecutor` with 4 threads.

The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)
```

`fit_model` is the function that we want to run asynchronously.

We instantiate a `ThreadPoolExecutor` with 4 threads.

A new job is submitted to the `Executor`.
The `Future` object `future1` holds the state of the computation.

The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)
```

`fit_model` is the function that we want to run asynchronously.

We instantiate a `ThreadPoolExecutor` with 4 threads.

A new job is submitted to the `Executor`.

The `Future` object `future1` holds the state of the computation.

The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...
```

`fit_model` is the function that we want to run asynchronously.

We instantiate a `ThreadPoolExecutor` with 4 threads.

A new job is submitted to the `Executor`.
The `Future` object `future1` holds the state of the computation.

The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...

    # Blocking call, wait and return the result
    model1 = future1.result(timeout=None)
    model2 = future2.result(timeout=None)
```

`fit_model` is the function that we want to run asynchronously.

We instantiate a `ThreadPoolExecutor` with 4 threads.

A new job is submitted to the `Executor`.
The `Future` object `future1` holds the state of the computation.

Wait for the computation to end and return the result with `f.result`.

The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...

    # Blocking call, wait and return the result
    model1 = future1.result(timeout=None)
    model2 = future2.result(timeout=None)

# The ressources have been cleaned up
print(model1, model2)
```

`fit_model` is the function that we want to run asynchronously.

We instantiate a `ThreadPoolExecutor` with 4 threads.

A new job is submitted to the `Executor`.
The `Future` object `future1` holds the state of the computation.

Wait for the computation to end and return the result with `f.result`.

The ressources are cleaned up.

The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...

    # Blocking call, wait and return the result
    model1 = future1.result(timeout=None)
    model2 = future2.result(timeout=None)

# The ressources have been cleaned up
print(model1, model2)
```

`fit_model` is the function that we want to run asynchronously.

We instantiate a `ThreadPoolExecutor` with 4 threads.

A new job is submitted to the `Executor`.
The `Future` object `future1` holds the state of the computation.

Wait for the computation to end and return the result with `f.result`.

The ressources are cleaned up.

Submitting more than one job returns an iterator: `executor.map`

The `Future` object: an asynchronous result state.

States

`Future` objects hold the state of the asynchronous computations, which can be in one of 4 states: `Not started`, `Running`, `Cancelled` and `Done`

The state of a `Future` can be checked using `f.running`, `f.cancelled`, `f.done`.

Blocking methods

- `f.result(timeout=None)`
- `f.exception(timeout=None)`

wait for computations to be done.

The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:


    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...

    # Blocking call, wait and return the result
    model1 = future1.result(timeout=None)
    model2 = future2.result(timeout=None)

# The resources have been cleaned up
print(model1, model2)
```



The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

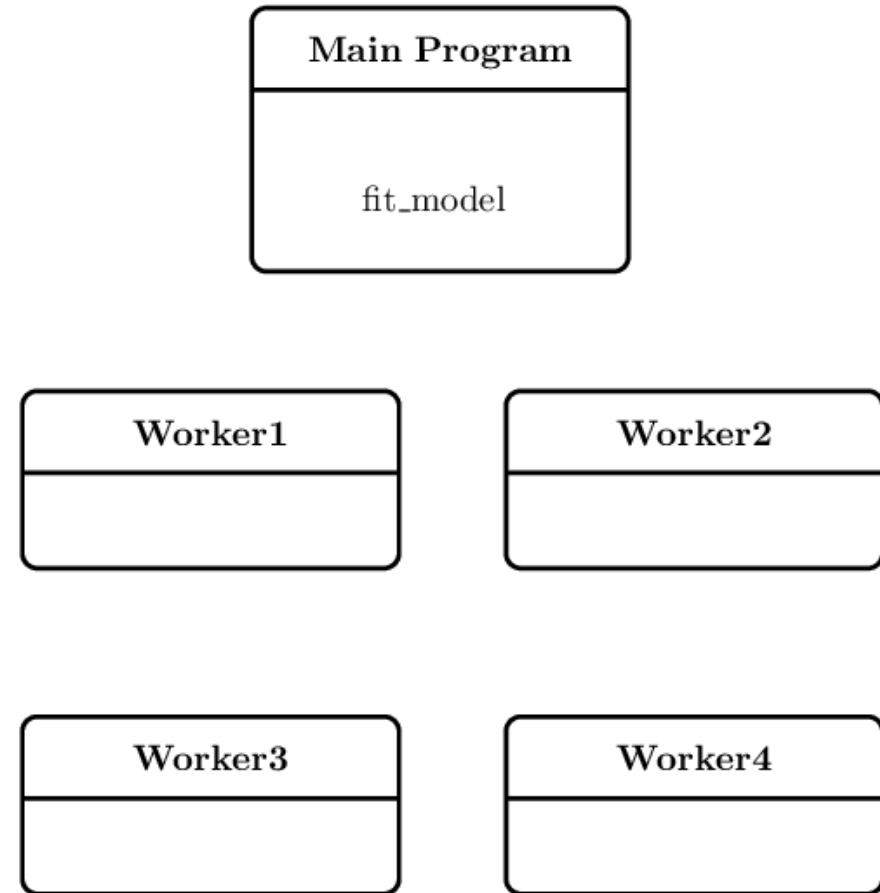
    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...

    # Blocking call, wait and return the result
    model1 = future1.result(timeout=None)
    model2 = future2.result(timeout=None)

# The ressources have been cleaned up
print(model1, model2)
```



The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

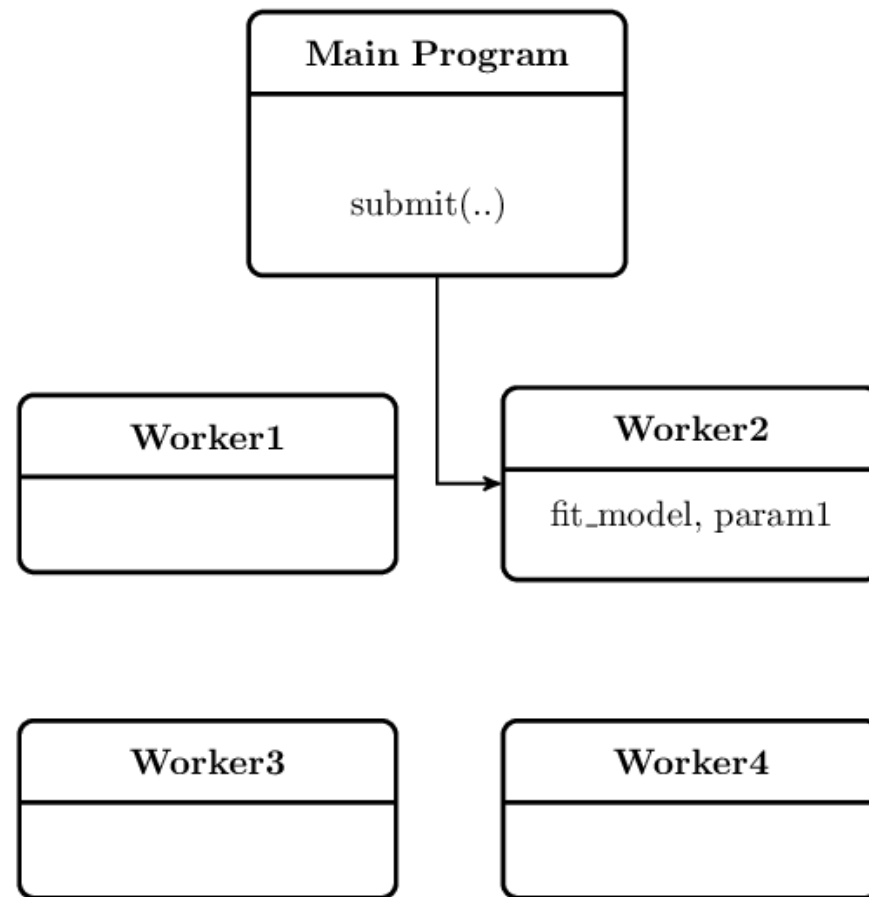
    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...

    # Blocking call, wait and return the result
    model1 = future1.result(timeout=None)
    model2 = future2.result(timeout=None)

# The resources have been cleaned up
print(model1, model2)
```



The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

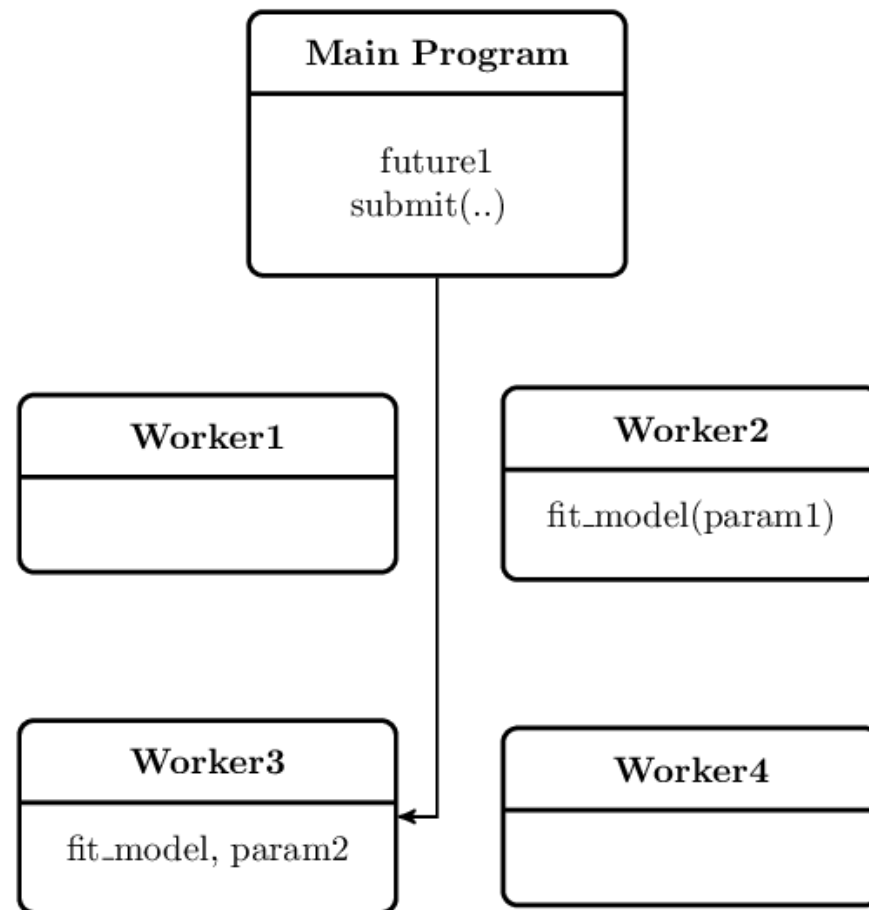
    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...

    # Blocking call, wait and return the result
    model1 = future1.result(timeout=None)
    model2 = future2.result(timeout=None)

# The resources have been cleaned up
print(model1, model2)
```



The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

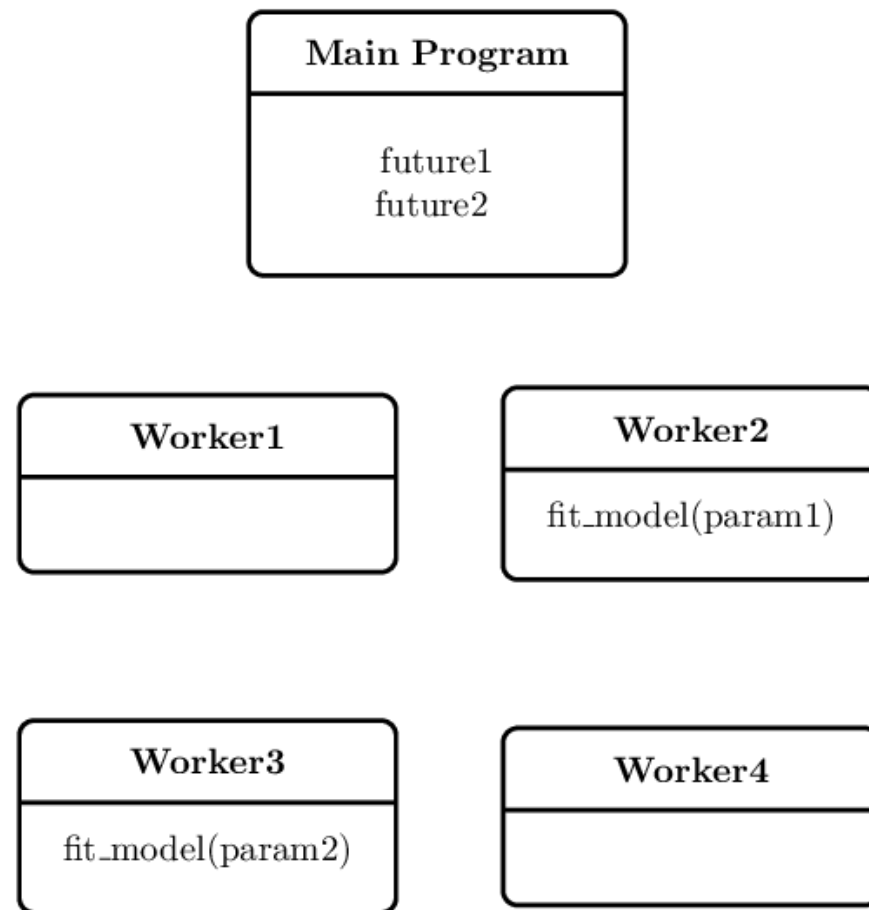
    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...

    # Blocking call, wait and return the result
    model1 = future1.result(timeout=None)
    model2 = future2.result(timeout=None)

# The ressources have been cleaned up
print(model1, model2)
```



The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

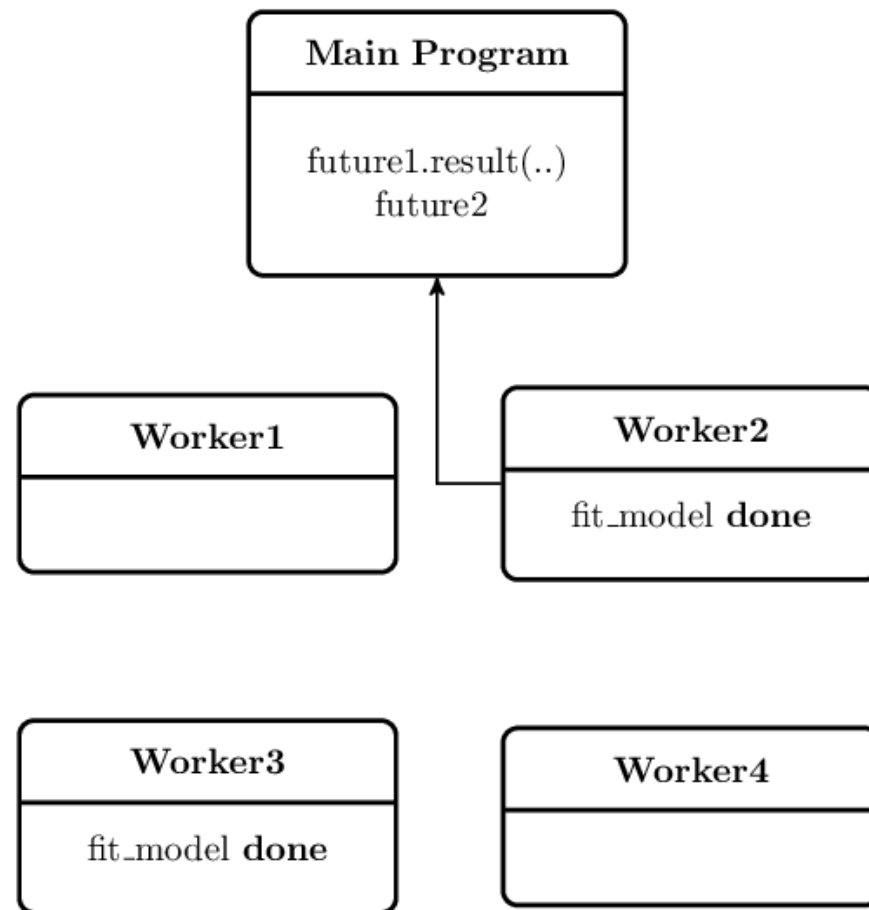
    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...

    # Blocking call, wait and return the result
    model1 = future1.result(timeout=None)
    model2 = future2.result(timeout=None)

# The resources have been cleaned up
print(model1, model2)
```



The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:

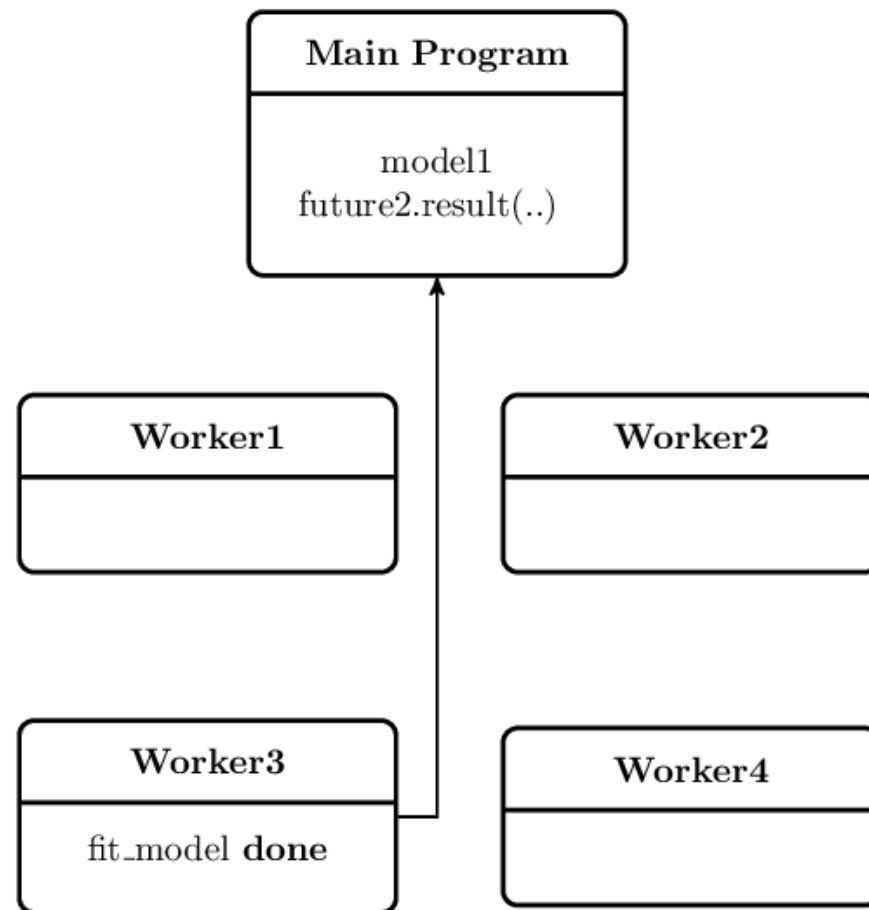
    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...

    # Blocking call, wait and return the result
    model1 = future1.result(timeout=None)
    model2 = future2.result(timeout=None)

# The resources have been cleaned up
print(model1, model2)
```



The Executor : a worker pool

```
from concurrent.futures import ThreadPoolExecutor

def fit_model(params):
    # Heavy computation
    return model

# Create an executor with 4 threads
with ThreadPoolExecutor(max_workers=4) as executor:


    # Submit an asynchronous job and return a Future
    future1 = executor.submit(fit_model, param1)

    # Submit other job
    future2 = executor.submit(fit_model, param2)

    # Run other computation
    ...

    # Blocking call, wait and return the result
    model1 = future1.result(timeout=None)
    model2 = future2.result(timeout=None)

# The ressources have been cleaned up
print(model1, model2)
```



Choosing the type of worker: Thread or
Process ?

Running on multiple cores

Python GIL

The internal implementation of python interpreter relies on a "Global Interpreter Lock" (**GIL**), protecting the concurrent access to python objects:

- Only one thread can acquire it.
- Not designed for efficient multicore computing.

Global lock everytime we access a python object.

Released when performing long I/O operations or by some libraries.
(*e.g.* numpy, openMP,..)

Thread

- Real system thread:
 - pthread
 - windows thread
- All the computation are done with a **single** interpreter.

Advantages:

- Fast spawning
- Reduced memory overhead
- No communication overhead (shared python objects)

Thread

- Real system thread:
 - pthread
 - windows thread
- All the computation are done with a **single** interpreter.

Advantages:

- Fast spawning
- Reduced memory overhead
- No communication overhead (shared python objects)

Wait... shared python objects and a single interpreter?!?

Thread

- Real system thread:
 - pthread
 - windows thread
- All the computation are done with a **single** interpreter.

Advantages:

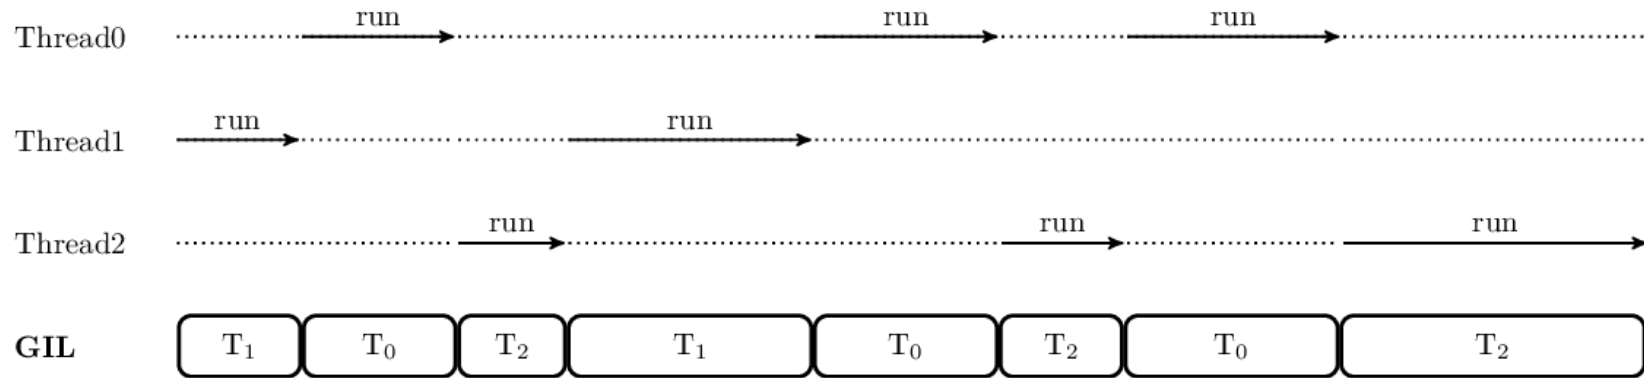
- Fast spawning
- Reduced memory overhead
- No communication overhead (shared python objects)

Wait... shared python objects and a single interpreter?!?

There is only one GIL!

Thread


Multiple threads running python code:




This is not quicker than sequential even on a multicore machine.

Thread

Threads hold the GIL when running python code.
They release it when blocking for I/O:



Or when using some c library:



Process

- Create a new python interpreter per worker.
- Each worker run in **its own** interpreter.

Inconvenients:

- *Slow* spawning
- Higher memory overhead
- Higher communication overhead.

Process

- Create a new python interpreter per worker.
- Each worker run in **its own** interpreter.

Inconvenients:

- *Slow* spawning
- Higher memory overhead
- Higher communication overhead.

But there is no GIL!

The computation can be done in parallel even for python code.

Process

- Create a new python interpreter per worker.
- Each worker run in **its own** interpreter.

Inconvenients:

- *Slow* spawning
- Higher memory overhead
- Higher communication overhead.

But there is no GIL!

The computation can be done in parallel even for python code.

Method to create a new interpreter: *fork* or *spawn*

Launching a new interpreter: *fork*

Duplicate the current interpreter. (Only available on UNIX)

Advantages:

- Low spawning overhead.
- The interpreter is warm *imported*.

Inconvenient:

- Bad interaction with multithreaded programs
- Does not respect the POSIX specifications

⇒ Some libraries crash: numpy on OSX, openMP, ...

Launching a new interpreter: *spawn*

Create a new interpreter from scratch.













Advantages:

- Safe (respect POSIX)
- Fresh interpreter without extra libraries.

Inconvenient:

- Slower to start.
- Need to reload the libraries, redefine the functions...

Comparison between Thread and Process

	Thread	Process (fork)	Process (spawn)
Efficient multicore run			
No communication overhead			
POSIX safe			
No spawning overhead			

Comparison between Thread and Process

	Thread	Process (fork)	Process (spawn)
Efficient multicore run	✗	✓	✓
No communication overhead	✓	✗	✗
POSIX safe	✓	✗	✓
No spawning overhead	✓	✓	Loky

⇒ Hide the spawning overhead by reusing the pool of processes.

Reusing a `ProcessPoolExecutor`.

Reusing a `ProcessPoolExecutor`.

To Avoid the spawning overhead, reuse a previously started `ProcessPoolExecutor`.

The spawning overhead is only paid once.

Easy using a global pool of process.

Main issue: is that robust?

Managing the state of the executor

Example deadlock

```
>>> from concurrent.futures import ProcessPoolExecutor
>>> with ProcessPoolExecutor(max_workers=4) as e:
...     e.submit(lambda: 1).result()
...
Traceback (most recent call last):
  File "/usr/lib/python3.6/multiprocessing/queues.py", line 241, in _feed
    obj = _ForkingPickler.dumps(obj)
  File "/usr/lib/python3.6/multiprocessing/reduction.py", line 51, in dumps
    cls(buf, protocol).dump(obj)
_pickle.PicklingError: Can't pickle <function <lambda> at 0x7fcff0184d08>:
attribute lookup <lambda> on __main__ failed

^C
```

It can be tricky to know which `submit` call crashed the `Executor`.

Managing the state of the executor

Example deadlock

```
>>> from concurrent.futures import ProcessPoolExecutor
>>> with ProcessPoolExecutor(max_workers=4) as e:
...     e.submit(lambda: 1)
...
Traceback (most recent call last):
  File "/usr/lib/python3.6/multiprocessing/queues.py", line 241, in _feed
    obj = _ForkingPickler.dumps(obj)
  File "/usr/lib/python3.6/multiprocessing/reduction.py", line 51, in dumps
    cls(buf, protocol).dump(obj)
_pickle.PicklingError: Can't pickle <function <lambda> at 0x7f5c787bd488>:
attribute lookup <lambda> on __main__ failed

^C
```

Even worse, shutdown itself is deadlocked.

Reusable pool of workers: `loky`.

loky : a robust pool of workers

```
>>> from loky import ProcessPoolExecutor
>>> class CrashAtPickle(object):
...     """Bad object that triggers a segfault at unpickling time."""
...     def __reduce__(self):
...         raise RuntimeError()
...
>>> with ProcessPoolExecutor(max_workers=4) as e:
...     e.submit(CrashAtPickle()).result()
...
Traceback (most recent call last):
...
RuntimeError
Traceback (most recent call last):
...
BrokenExecutor: The QueueFeederThread was terminated abruptly while feeding a
new job. This can be due to a job pickling error.
>>>
```

- Return and raise a user friendly exception.
- Fix some other deadlocks.

A reusable ProcessPoolExecutor

```
>>> from loky import get_reusable_executor
>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
0
```

```
>>> excutor.submit(id, 42).result()
139655595838272
```

```
>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
0
```

```
>>> excutor.submit(CrashAtUnpickle()).result()
Traceback (most recent call last):
...
BrokenExecutorError
>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
1
>>> excutor.submit(id, 42).result()
139655595838272
```

Create a `ProcessPoolExecutor` using the factory function `get_reusable_executor`.

A reusable ProcessPoolExecutor

```
>>> from loky import get_reusable_executor
>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
0

>>> excutor.submit(id, 42).result()
139655595838272

>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
0

>>> excutor.submit(CrashAtUnpickle()).result()
Traceback (most recent call last):
...
BrokenExecutorError
>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
1
>>> excutor.submit(id, 42).result()
139655595838272
```

Create a `ProcessPoolExecutor` using the factory function `get_reusable_executor`.

The executor can be used exactly as `ProcessPoolExecutor`.

A reusable ProcessPoolExecutor

```
>>> from loky import get_reusable_executor
>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
0

>>> excutor.submit(id, 42).result()
139655595838272

>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
0

>>> excutor.submit(CrashAtUnpickle()).result()
Traceback (most recent call last):
...
BrokenExecutorError
>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
1
>>> excutor.submit(id, 42).result()
139655595838272
```

Create a `ProcessPoolExecutor` using the factory function `get_reusable_executor`.

The executor can be used exactly as `ProcessPoolExecutor`.

When the factory is called elsewhere, reuse the same executor if it is working.

A reusable `ProcessPoolExecutor`

```
>>> from loky import get_reusable_executor
>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
0

>>> excutor.submit(id, 42).result()
139655595838272

>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
0

>>> excutor.submit(CrashAtUnpickle()).result()
Traceback (most recent call last):
...
BrokenExecutorError
>>> excutor = get_reusable_executor(max_workers=4)
>>> print(excutor.executor_id)
1

>>> excutor.submit(id, 42).result()
139655595838272
```

Create a `ProcessPoolExecutor` using the factory function `get_reusable_executor`.

The executor can be used exactly as `ProcessPoolExecutor`.

When the factory is called elsewhere, reuse the same executor if it is working.

When the executor is broken, automatically re-spawn a new one.

Conclusion

Conclusion

- `Thread` can be efficient to run multicore programs if your code releases the **GIL**.
- Else, you should use `Process` with `spawn` and try to reuse the pool of process as much as possible.
- `loky` can help you do that ;).
- Improves the management of a pool of workers in projects such as `joblib`.

Thanks for your attention!

Slides available at tommoral.github.io/pyparis17/

More on the GIL by Dave Beazley : dabeaz.com/python/GIL.pdf

🐙 Loky project : github.com/tommoral/loky

🐦 [@tomamoral](https://twitter.com/tomamoral)