

Kickoff ExaDost, Paris 2023.09.13

Scaling up Scikit-Learn for Exascale



MIND



Scikit-learn:



The reference library for ML on tabular data

- Many reference algorithms (no deep learning)
- Tools designed to create ML Pipelines (preprocessing)
- Utilities to perform hyperparameter optimization
- Tools to evaluate the results

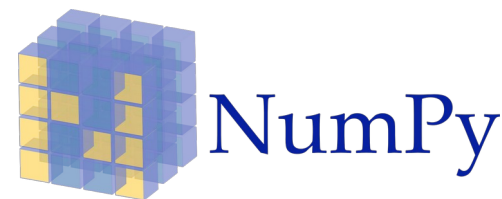
Scikit-learn:

The reference library for ML on tabular data



Current constraints

- Input are pandas or numpy arrays
- Parallelism is done using joblib
- No GPU acceleration (for now)



A zoom in the future development of Scikit-Learn

I - Parallelism in scikit-learn

II – The Array API

III - Exploring a plugin system in scikit-learn

A high level API with joblib:

```
>>> from math import sqrt
>>> [sqrt(i ** 2) for i in range(10)]
[0.0, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0]
```

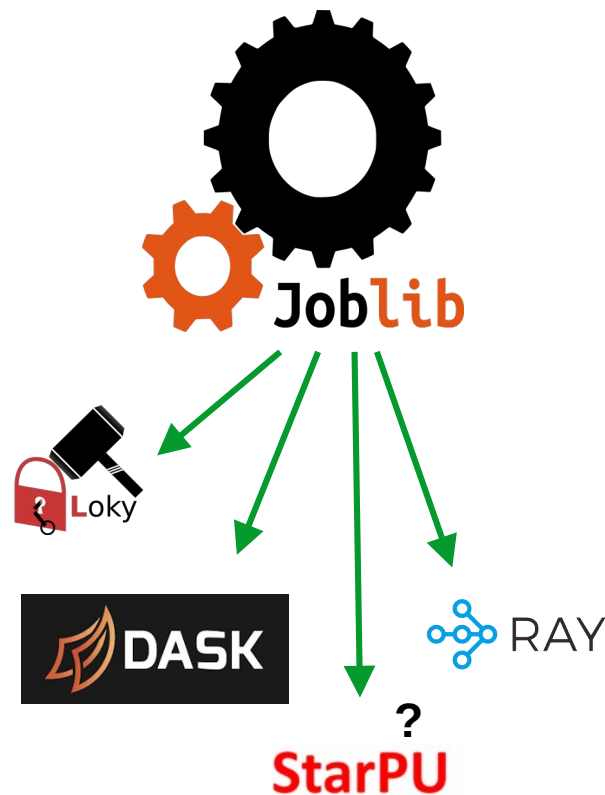


```
>>> from math import sqrt
>>> from joblib import Parallel, delayed
>>> Parallel(n_jobs=2)(delayed(sqrt)(i ** 2) for i in range(10))
[0.0, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0]
```



```
>>> from joblib import parallel_config
>>> with parallel_config(backend='threading', n_jobs=2):
...     Parallel()(delayed(sqrt)(i ** 2) for i in range(10))
[0.0, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0]
```

> Embarassingly Parallel tasks, with pluggable backends



Low-level parallelism:

Calls to BLAS (through numpy) are parallelised

Low level parallelism using openMP



> Intrinsic parallelism, dependent on the data structure and the algorithm.

Challenge: sometime hard to make the different parallelism levels coexist!

A zoom in the future development of Scikit-Learn

I - Parallelism in scikit-learn

II – The Array API

III - Exploring a plugin system in scikit-learn

Introduction to the array API standard:

```
1 import numpy as np
2 data = np.arange(9).reshape(3, 3)
3 # with numpy, only get one output
4 # uses keyword `axis`
5 max_over_cols = np.max(data, axis=1)
```


Introduction to the array API standard:

```
1 import numpy as np
2 data = np.arange(9).reshape(3, 3)
3 # with numpy, only get one output
4 # uses keyword `axis`
5 max_over_cols = np.max(data, axis=1)
```

```
1 import torch
2 data = torch.arange(9).reshape(3, 3)
3 # with pytorch, `torch.max` on nd tensors
4 # returns a length 2 tuple
5 # `max` expects parameter `dim`
6 max_over_cols, _ = torch.max(data, dim=1)
```

Introduction to the array API standard:

```
1 import numpy as np
2 data = np.arange(9).reshape(3, 3)
3 # with numpy, only get one output
4 # uses keyword `axis`
5 max_over_cols = np.max(data, axis=1)
```

```
1 import torch
2 data = torch.arange(9).reshape(3, 3)
3 # with pytorch, `torch.max` on nd tensors
4 # returns a length 2 tuple
5 # `max` expects parameter `dim`
6 max_over_cols, _ = torch.max(data, dim=1)
```

Different array libraries can have different syntax for the same operation !

Introduction to the array API standard:

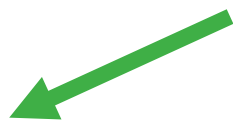
The Array API enables writing **agnostic single-source code** and interface it with all arrays !

```
1  from array_api_compat import get_namespace
2  def max_over_cols_fn(data):
3      xp = get_namespace(data)
4      return xp.max(data, axis=1)
```

Introduction to the array API standard:

The Array API enables writing **agnostic single-source code** and interface it with all arrays !

```
1  from array_api_compat import get_namespace
2  def max_over_cols_fn(data):
3      xp = get_namespace(data)
4      return xp.max(data, axis=1)
```



```
1  import numpy as np
2  data = np.arange(9).reshape(3, 3)
3  max_over_cols = max_over_cols_fn(data)
```



```
1  import torch
2  data = torch.arange(9).reshape(3, 3)
3  max_over_cols = max_over_cols_fn(data)
```

Array API spec excerpt:

Assumptions

API

Design topics & constraints

Future API standard evolution

API specification

Array object

Broadcasting

Constants

Creation Functions

Data Type Functions

Data Types

Element-wise Functions

Function and method
signatures

Indexing

Indexing Functions

Linear Algebra Functions

Manipulation Functions

Searching Functions

A conforming implementation of the array API standard must provide and support the

Objects in API

<code>broadcast_arrays(*arrays)</code>	Broadcasts one or more arrays to a common shape
<code>broadcast_to(x, /, shape)</code>	Broadcasts an array to a given shape
<code>concat(arrays, /, *, axis=0)</code>	Joins a sequence of arrays along an axis
<code>expand_dims(x, /, *, axis=0)</code>	Expands the shape of an array by inserting a new axis of size one at the position of the axis parameter
<code>flip(x, /, *, axis=None)</code>	Reverses the order of the elements in the array along the axis or axes given by axis
<code>permute_dims(x, /, axes)</code>	Permutes the axes of an array according to the order given by axes
<code>reshape(x, /, shape, *, copy=None)</code>	Reshapes an array without changing its data
<code>roll(x, /, shift, *, axis=None)</code>	Rolls array elements along a given axis
<code>squeeze(x, /, axis)</code>	Removes singleton dimensions from an array
<code>stack(arrays, /, *, axis=0)</code>	Joins a sequence of arrays along a new axis

https://data-apis.org/array-api/latest/API_specification/index.html

Compatible array libraries:

(done or ongoing)
(non-exhaustive)



CPU
DISTRIBUTED



GPU NVIDIA

NumPy 
CPU

 **PyTorch**

CPU
GPU AMD/INTEL/NVIDIA
OTHERS

Dpctl + DPNP

CPU
GPU INTEL

Toward Array API compliancy in scikit-learn:

When possible, use the Array Api rather than numpy calls !

```
636 - np.square(X, out=X)
637 - np.sum(X, axis=0, out=X[0])
638 - total_var = (X[0] / N).sum()
639
640     self.explained_variance_ratio_ =
        self.explained_variance_ / total_var
641 - self.singular_values_ = S.copy() # Store the singular
        values.
642
643     if self.n_components_ < min(n_features, n_samples):
644 - self.noise_variance_ = total_var -
        self.explained_variance_.sum()
```

```
659 + X **= 2
660 + total_var = xp.sum(xp.sum(X, axis=0) / N)
661
662     self.explained_variance_ratio_ =
        self.explained_variance_ / total_var
663 + self.singular_values_ = xp.asarray(S, copy=True) #
        Store the singular values.
664
665     if self.n_components_ < min(n_features, n_samples):
666 + self.noise_variance_ = total_var -
        xp.sum(self.explained_variance_)
```

Excerpt of



<https://github.com/scikit-learn/scikit-learn/pull/26315>

ENH Array API support for PCA

Towards Array API compliancy in scikit-learn:



[sklearn.lda.LDA](#)



[sklearn.decomposition.PCA](#)

(next release)



[sklearn.preprocessing.MinMaxScaler](#)

(ongoing)



Discussions and roadmaps [#22352](#) [#26024](#) [#26083](#)

Towards Array API compliancy in scikit-learn:

🔥 *How to use it ?* 🔥

```
1 from sklearn import config_context
2 from sklearn.decomposition import PCA
3 pca_estimator = PCA(n_components=10)
4 with config_context(array_api_dispatch=True):
5     pca_estimator.fit(X)
```

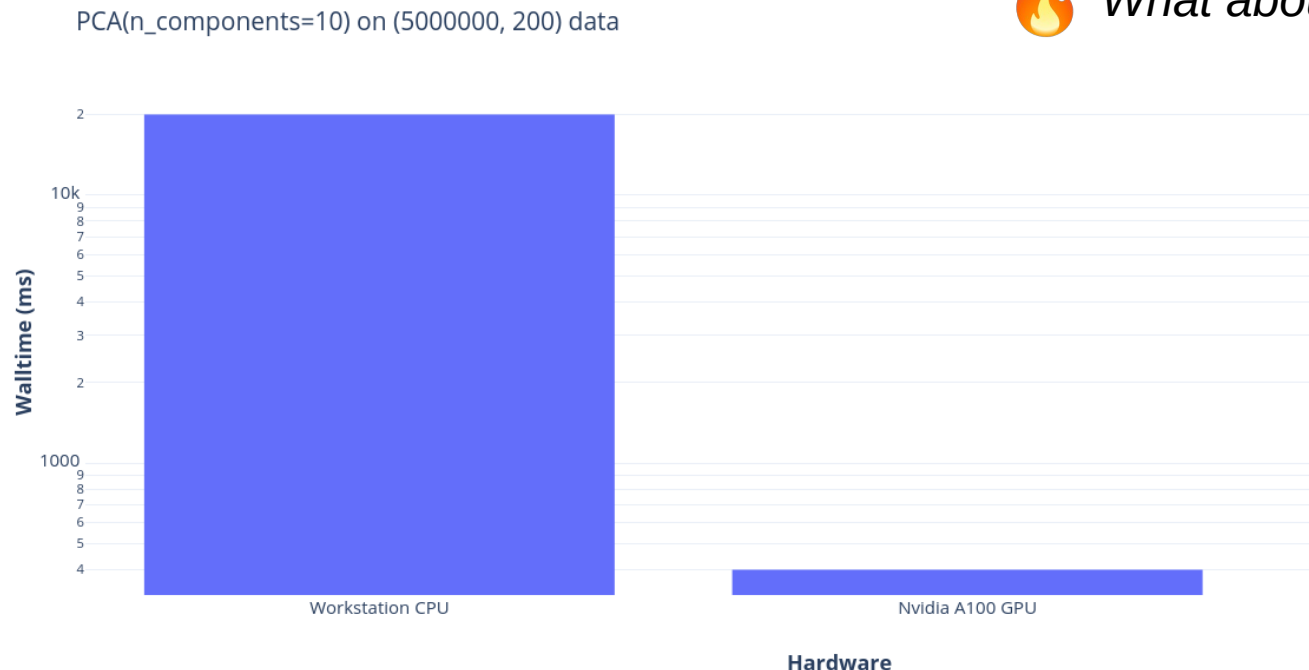


Experimental features require explicit activation in scikit-learn config class

```
1 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
2 lda = LinearDiscriminantAnalysis()
3 with config_context(array_api_dispatch=True):
4     lda.fit(X)
```

Towards Array API compliancy in scikit-learn:


🔥 *What about performance ?* 🔥



YES when CPU takes too long
40x speedup with pytorch
on GPU




NB: pytorch backend for intel
GPUs currently lacks optimization
for QR factorization, so Intel GPUs
are not included in the figure

Towards Array API compliancy in scikit-learn:

 *Will all estimators support all compatible array libraries ?* 

Likely not :-)

Current limitations of the Array API approach:

-  Does not cover all functions for all libraries (torch.topk, np.argpartition, RNG...)
-  Lacks synchronization functions for libraries that evaluate lazily (dask, jax...)
-  Not all algorithms can be implemented (efficiently) with high level array operations !

(random forest, histogram gradient boosting trees,...)

Possible drawbacks for the Array API:

Each call to a high-level array function writes the result into memory.

Example of pairwise_distance + argmin common pattern:

```
1  import numpy as np
2  from scipy.spatial.distance import cdist
3
4  # create some data
5  rng = np.random.default_rng(123)
6  data = rng.random((500000, 200))
7  query = rng.random((1, 200))
8
9  # writes all pairwise distances to memory
10 pairwise_distances = cdist(query, data)
11
12 closest = np.argmin(pairwise_distances)
```

A zoom in the future development of Scikit-Learn

I - Parallelism in scikit-learn

II – The Array API

III - Exploring a plugin system in scikit-learn

Exploring a plugin system in scikit-learn:

> Preview of user workflow:

```
$  
$ pip install scikit-learn some-compute-plugin
```

Exploring a plugin system in scikit-learn:

> Preview of user workflow:

```
$  
$ pip install scikit-learn some-compute-plugin
```

```
1 from sklearn import config_context  
2 from sklearn.cluster import KMeans  
3 kmeans_estimator = KMeans()  
4 with config_context(engine_provider="some_compute_plugin"):  
5     y_pred = kmeans_estimator.fit_predict(X)  
6     y_transform = kmeans_estimator.transform(X)
```

Exploring a plugin system in scikit-learn:

> Example for plugin developers: register a custom engine class

```
1 class MyKMeansEngine:
2
3     def __init__(self, estimator):
4         self.estimator = estimator
5
6     def prepare_fit(self, X, y=None, sample_weight=None):
7         """TODO: insert custom engine implementation"""
8
9     def kmeans_single(self, X, sample_weight, centers_init):
10        """TODO: insert custom engine implementation"""
11
12    def prepare_prediction(self, X, sample_weight):
13        """TODO: insert custom engine implementation"""
14
15    def get_labels(self, X, sample_weight):
16        """TODO: insert custom engine implementation"""
```

> Test against native scikit-learn unit tests

```
~# pytest --sklearn-engine-provider some-compute-plugin --pyargs sklearn.cluster.tests.test_k_means
```


Exploring a plugin system in scikit-learn:



Foundational issue

<https://github.com/scikit-learn/scikit-learn/issues/22438>



Pull request

<https://github.com/scikit-learn/scikit-learn/pull/25535>



Development branch

<https://github.com/scikit-learn/scikit-learn/tree/feature/engine-api>



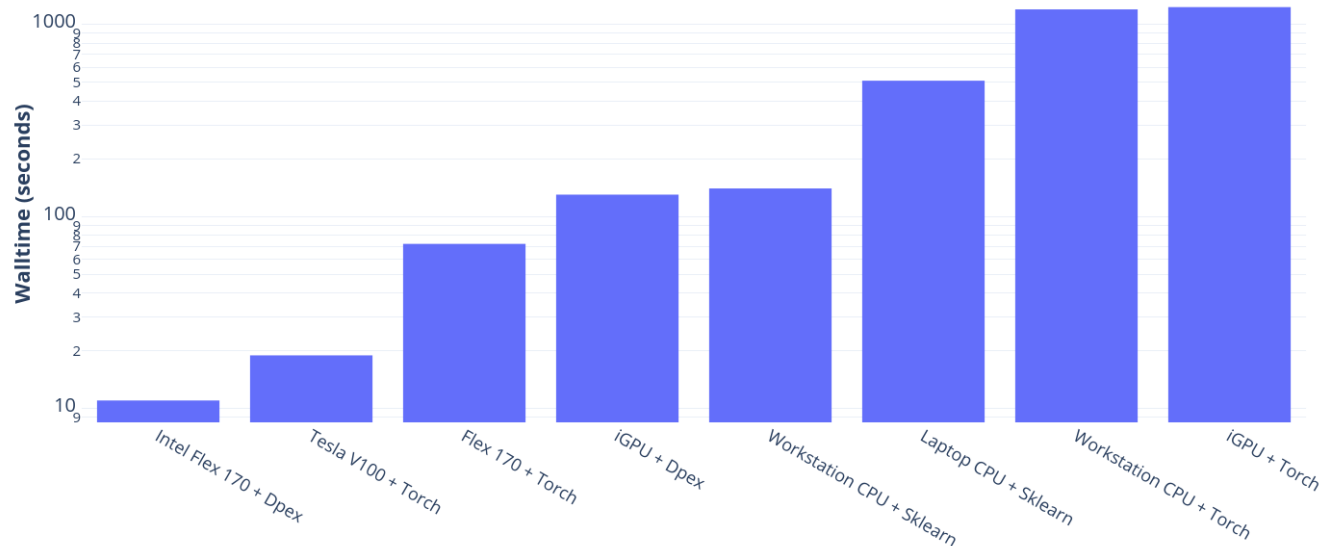
Example plugin for KMeans on intel GPUs

<https://github.com/soda-inria/sklearn-numba-dpex>

A plugin for k-means with the ¹oneAPI toolchain:

oneAPI

Walltime for 100 k-means lloyd iterations on 50_000_000 * 14 data



k-means is one order of magnitude faster on GPU, more significantly so when optimized with a low-level implementation

NB: cuml implementation of k-means for Nvidia GPUs could not be included in the figure because it could not scale to this amount of data memory-wise.

sklearn_numba_dpex

device + software backend



<https://github.com/soda-inria/sklearn-numba-dpex>

ROADMAP in ExaDost:

- 1 Engineer position to work on scikit-learn and joblib (MIND)
 - Work on the lazy Array API
 - Work on online computations with scikit-learn (partial_fit),
 - Work on improving nested parallelism handling.
- 1 PhD position with H. Hendrikx, (Thot)
 - Improve the distributed algorithms in scikit-learn